

Linear Regression



Disclaimer: Some content and images are borrowed from online sources listed in the reference section

Marks Distribution: Updated



•	Mid 1	15 %
•	Mid 2	15 %
•	Assignments/project	35 %
•	Daily Assessment	05 %
•	End-semester exam	30 %

Problem Formulation

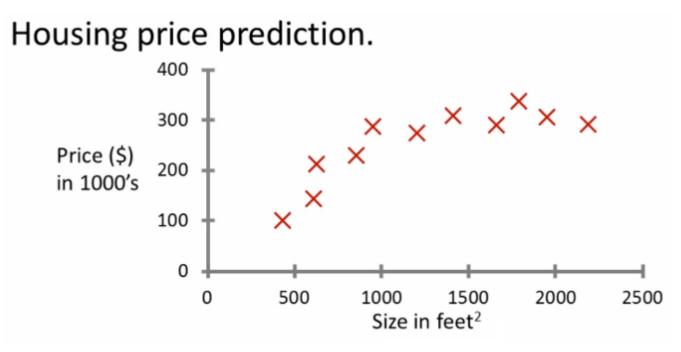


• Given a the training set of pairs $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), ..., (x^{(m)}, y^{(m)})$, where $x^{(i)} \in \mathbb{R}^d$ and $y^{(i)}$ is a **continuous** target variable, the task is to predict for $x^{(m+j)}, j \geq 1$.

Let us start with an Example



- How do we predict housing prices
 - Collect data regarding housing prices and how they relate to size in feet.



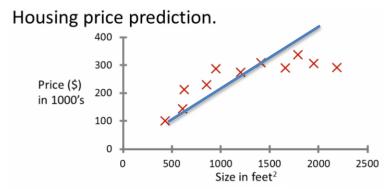
Example problem



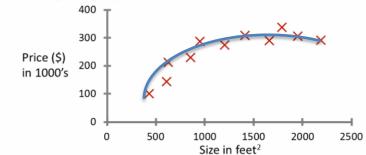
- "Given this data, a friend has a house 750 square feet
 - how much can they be expected to get?"

- Straight line through data
 - Maybe \$150 000

- Second order polynomial
 - Maybe \$200 000



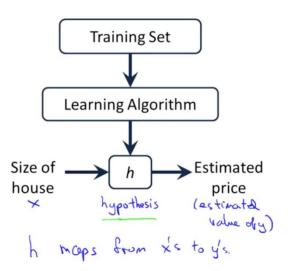




Linear Regression



- With our training set defined how do we use it?
 - Take training set
 - Pass into a learning algorithm
 - Algorithm outputs a function (denoted h)
 - This function takes an input (e.g. size of new house)
 - Tries to output the estimated value of Y
- How do we represent hypothesis h?
 - Going to present h as $h_{\theta}(x) = \theta_0 + \theta_1 x$
 - Means Y is a linear function of x
 - $-\theta_i$ are **parameters**

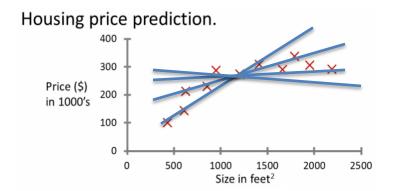


- A linear regression with one variable is also called univariate linear regression
- So in summary
 - A hypothesis takes in some variable
 - Uses parameters determined by a learning system
 - Outputs a prediction based on that input

Which line is best?



- Many lines are possible !!
 Which is the best?
- A cost function lets us figure out how to fit the best straight line to our data.



- What makes a line different?
 - Parameters θ_0 , θ_1
- Which is the best line?
 - The line that minimizes the difference between the actual and estimated prices.
- What is our objective?
 - Choose these parameters θ_0 , θ_1 so that $h_{\theta}(x)$ is close to y for our training examples, i.e. minimize the difference between h(x) and y for each/any/every example.

Objective

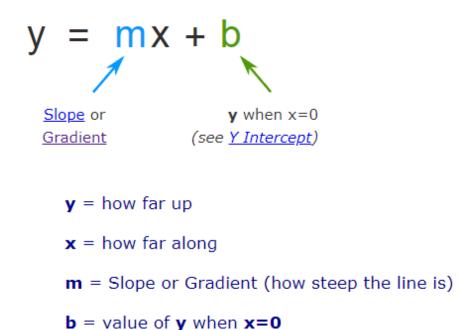


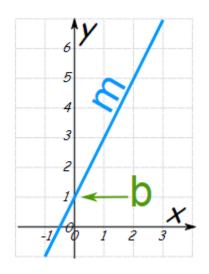
• Loss function:
$$J(\theta) = \frac{1}{2m} \sum_{i=1}^{m} (h_{\theta}(x^i) - y^i)^2$$

- Minimize $J(\theta)$ θ_0, θ_1
- How do we achieve this?
 - That is where gradient descent helps us!

Gradient



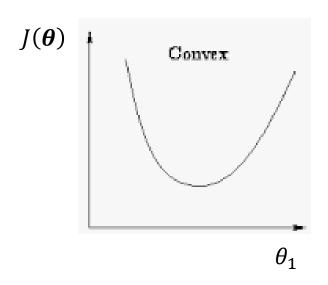


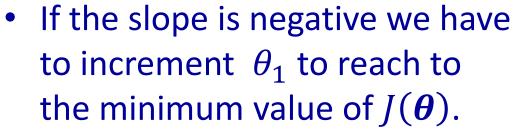


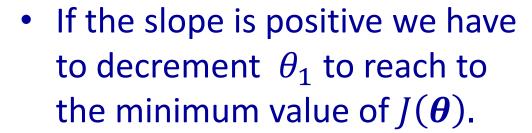


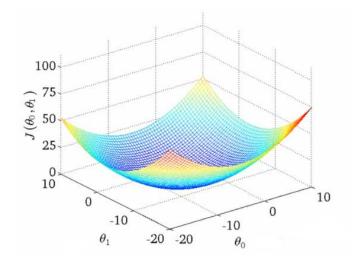
The gradient is just the vector of <u>partial derivatives</u>

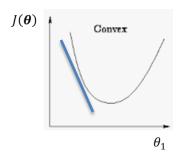
Graphical representation of a convex cost function

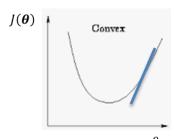










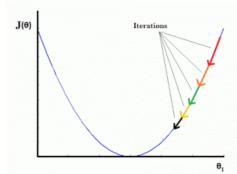


Is this okay?

Initialize θ_0 , θ_1

Repeat the following until convergence

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad \text{(for } j = 0 \text{ and } j = 1)$$



Gradient
Descent
Algorithm

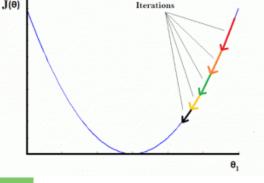
– That is when the gradient is negative we are incrementing θ_j and when the gradient is positive we are decrementing θ_j .

• j=0:
$$\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^i) - y^i)$$

• j=1:
$$\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^i) - y^i) \cdot x^i$$

Gradient Descent

So, Gradient descent is an optimization algorithm used to find the values of parameters of a function that minimizes a cost function (cost).



 Gradient descent is best used when the parameters cannot be calculated analytically (e.g. using linear algebra) and must be searched for by an optimization algorithm.

Multivariate Linear Regression



Linear Regression with multiple input variables/features.

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_1	x_2	x_3	x_4	y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178

More notations

- n: number of features (n = 4)
- m: number of examples (i.e. number of rows in a table)
- $-\mathbf{x}^{i}$: vector of the input for an example (so a vector of the four parameters for the ith input example)
 - x³ is, for example, the 3rd house, and contains the four features associated with that house
- x_jⁱ The value of feature j in the ith training example
 x₂³ is, for example, the number of bedrooms in the third house

What is the form of our hypothesis?



- Previously our hypothesis took the following form.
 - $-h_{\theta}(x) = \theta_0 + \theta_1 x$
 - Here we have two parameters (theta 1 and theta 2) determined by our cost function
 - One variable x
- Now we have multiple features

$$- h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$$

• Let us take $x_0 = 1$ for convenience of notation

$$- h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$$

$$-h_{\theta}(x) = \theta^T X$$

Gradient descent for multiple variables



Our cost function is

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient descent

Repeat
$$\{$$
 $o \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$ $\}$ (simultaneously update for every $j=0,\dots,n$)

Gradient decent for multiple variables



When n = 1

Repeat
$$\left\{ \begin{array}{l} \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \\ \hline \frac{\partial}{\partial \theta_0} J(\theta) \end{array} \right.$$
 $\left. \theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)} \right.$ (simultaneously update θ_0, θ_1) $\left. \right\}$

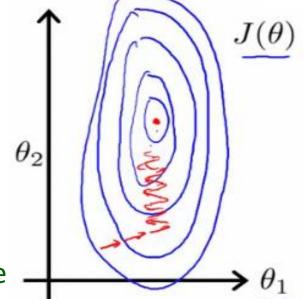
• When n >=1

New algorithm
$$(n \ge 1)$$
: Repeat $\left\{ \begin{array}{c} \frac{2}{\sqrt{2}} & \frac{2}{\sqrt{2}} & \frac{2}{\sqrt{2}} \\ \theta_j := \theta_j - \alpha & \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} \\ & \text{(simultaneously update } \theta_j \text{ for } \\ & j = 0, \dots, n) \end{array} \right\}$

Gradient Decent in practice: 1 Feature Scaling

Mahindra École Centrale COLLEGE OF ENGINEERING

- Range difference of features
 - x1 = size (0 2000 feet)
 - x2 = number of bedrooms (1-5)
 - Means the contours generated if we plot θ_1 vs. θ_2 give a very tall and thin shape due to the huge range difference



- Running gradient descent on this kind of cost function can take a long time to find the global minimum
- Feature scaling
 - If you have a problem with multiple features
 - You should make sure those features have a similar scale
 - Means gradient descent will converge more quickly

Some feature scaling methods

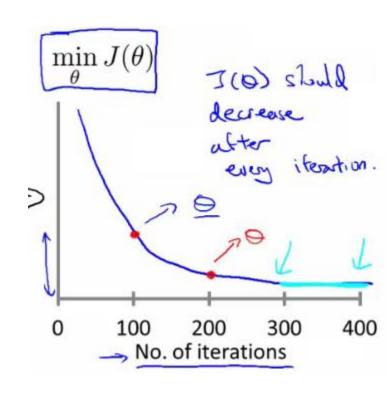


- Max/Min Scaling
- Mean Normalization
- Z-score Scaling

Learning Rate α



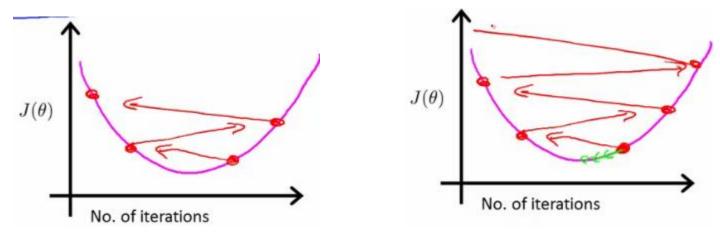
- Focus on the learning rate (α)
 - How to chose α ?
- Make sure gradient descent is working
- Plot min $J(\theta)$ vs. no of iterations
 - (i.e. plotting $J(\theta)$ over the course of gradient descent
- If gradient descent is working then J(θ) should decrease after every iteration



Learning Rate α



- Checking its working If you plot $J(\theta)$ vs. iterations and see the value is increasing means you probably need a smaller α
 - Cause is because your minimizing a function which looks like this



 But you overshoot, so reduce learning rate so you actually reach the minimum (green line)

•

Vector Notation

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_1	x_2	x_3	x_4	y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178

•
$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}_{m \times (n+1)} Y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}_{m \times 1}$$

•
$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}_{(n+1) \times 1}$$
 $h_{\theta}(x) = x_0 \theta_0 + \dots + x_n \theta_n = X\theta$

•
$$J = \frac{1}{2m} (X\theta - Y)^T_{1 \times m} (X\theta - Y)_{m \times 1}$$

Vectorization

New algorithm
$$(n \ge 1)$$
:

Repeat $\left\{ \begin{array}{c} \frac{\partial}{\partial x_j} \mathcal{T}(\mathbf{S}) \\ \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \\ \text{(simultaneously update } \theta_j \text{ for } j = 0, \dots, n) \right\}$

•
$$\frac{\partial}{\partial \theta_i} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^i) - y^i) \cdot x_j^i$$

$$= \frac{1}{m} \sum_{i=1}^{m} ((x^{i}\theta) - y^{i}) \cdot x_{j}^{i}$$

•
$$\Theta_{(n+1)\times 1} = \Theta_{(n+1)\times 1} - \frac{\alpha}{m} (X^T_{(n+1)\times m} (X\theta - Y)_{m\times 1})$$

Closed form solution



• Given X and Y, our aim is to find θ so that $Y = X\theta$.

- $Y = X\theta$.
- $X^{-1}Y = \theta$
- $Y^{-1}(X^{T^{-1}}X^T)Y = \theta$
- $(X^TX)^{-1}X^TY = \theta$

Why not closed form?



- The problem with this operation is the time complexity of calculating the inverse of a nxn matrix which is O(n^3) and as n increases it can take a very long time to finish.
- When n is low (n < 1000 or n < 10000) you can think
 of normal equations as the better option for
 calculation theta, however for greater
 values Gradient Descent is much more faster, so the
 only reason is the time.

• Closed form works better when the input size is smaller. No need to choose α and no need to iterate.

Batch, Stochastic and Mini-Batch | Gradient Descent

 In Gradient Descent or Batch Gradient Descent, we use the whole training data for updating theta (or, w).

$$w = w - \alpha \nabla_w J(w) \tag{6}$$

• In Stochastic Gradient Descent, we use only single training example for updating theta.

$$w = w - \alpha \nabla_w J(x^i, y^i; w) \tag{7}$$

 Mini-batch Gradient Descent lies in between of these two extremes, in which we can use a mini-batch(small portion) of training data for updating theta.

$$w = w - \alpha \nabla_w J(x^{\{i:i+b\}}, y^{\{i:i+b\}}; w)$$
(7)

References



- Bishop, Christopher M. "Pattern recognition and machine learning, 2006." Spinger 60.1 (2012): 78-78.
- 2) http://www.holehouse.org/mlclass/
- 3) <u>https://www.mathsisfun.com/equation_of_line.htm</u> <u>l</u>
- 4) https://towardsdatascience.com/gradient-descent-algorithm-and-its-variants-10f652806a3