

# Linear Regression

*Venkateswara*

Disclaimer: Some content and images are borrowed from online sources listed in the reference section

# Marks Distribution: Updated

• Mid 1	15 %
• Mid 2	15 %
• Assignments/project	35 %
• Daily Assessment	05 %
• End-semester exam	30 %


# Problem Formulation

- Given a the training set of pairs  $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$ , where  $x^{(i)} \in \mathbb{R}^d$  and  $y^{(i)}$  is a **continuous** target variable, the task is to predict for  $x^{(m+j)}, j \geq 1$ .

# Let us start with an Example

- How do we predict housing prices
  - Collect data regarding housing prices and how they relate to size in feet.


Housing price prediction.




# Example problem

- "Given this data, a friend has a house 750 square feet - how much can they be expected to get?"
- Straight line through data
  - Maybe \$150 000
- Second order polynomial
  - Maybe \$200 000

Housing price prediction.

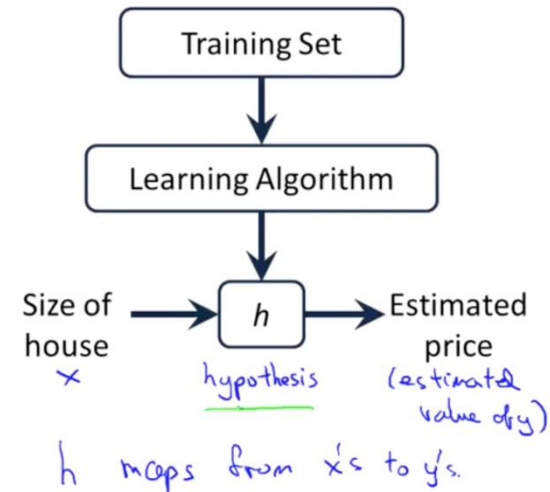


Housing price prediction.



# Linear Regression


- With our training set defined - how do we use it?
  - Take training set
  - Pass into a learning algorithm
  - Algorithm outputs a function (denoted  $h$ )
  - This function takes an input (e.g. size of new house)
    - Tries to output the estimated value of  $Y$
- How do we represent hypothesis  $h$  ?
  - Going to present  $h$  as  $h_{\theta}(x) = \theta_0 + \theta_1 x$
  - Means  $Y$  is a linear function of  $x$
  - $\theta_i$  are **parameters**
- A linear regression with one variable is also called **univariate linear regression**
- So in summary
  - A hypothesis takes in some variable
  - Uses parameters determined by a learning system
  - Outputs a prediction based on that input



# Which line is best ?

- Many lines are possible !!  
Which is the best?
- A cost function lets us figure out how to fit the best straight line to our data.
- What makes a line different ?
  - Parameters  $\theta_0, \theta_1$
- Which is the best line ?
  - The line that minimizes the difference between the actual and estimated prices.
- What is our objective ?
  - Choose these parameters  $\theta_0, \theta_1$  so that  $h_{\theta}(x)$  is close to  $y$  for our training examples, i.e. minimize the difference between  $h(x)$  and  $y$  for each/any/every example.

Housing price prediction.



# Objective

- Loss function:  $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$
- *Minimize  $J(\theta)$*   
 $\theta_0, \theta_1$
- How do we achieve this ?
  - That is where gradient descent helps us !




# Gradient

$$y = mX + b$$

Slope or  
Gradient

$y$  when  $x=0$   
(see Y Intercept)



$y$  = how far up

$x$  = how far along




$m$  = Slope or Gradient (how steep the line is)

$b$  = value of  $y$  when  $x=0$





- The gradient is just the vector of partial derivatives

# Graphical representation of a convex cost function



- If the slope is negative we have to increment  $\theta_1$  to reach to the minimum value of  $J(\theta)$ .
- If the slope is positive we have to decrement  $\theta_1$  to reach to the minimum value of  $J(\theta)$ .




# Is this okay ?

Initialize  $\theta_0, \theta_1$

- Repeat the following until convergence

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$

- That is when the gradient is negative we are incrementing  $\theta_j$  and when the gradient is positive we are decrementing  $\theta_j$ .




Gradient  
Descent  
Algorithm

- $j=0: \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)$
- $j=1: \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) \cdot x^i$

# Gradient Descent

- So, Gradient descent is an optimization algorithm used to find the values of parameters of a function that minimizes a cost function (cost).
- Gradient descent is best used when the parameters cannot be calculated analytically (e.g. using linear algebra) and must be searched for by an optimization algorithm.



# Multivariate Linear Regression

- Linear Regression with multiple input variables/features.

Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
$x_1$	$x_2$	$x_3$	$x_4$	$y$
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178

- More notations
  - $n$ : number of features ( $n = 4$ )
  - $m$ : number of examples (i.e. number of rows in a table)
  - $\mathbf{x}^i$ : vector of the input for an example (so a vector of the four parameters for the  $i^{\text{th}}$  input example)
    - $\mathbf{x}^3$  is, for example, the 3rd house, and contains the four features associated with that house
  - $x_j^i$ : The value of feature  $j$  in the  $i^{\text{th}}$  training example
    - $x_2^3$  is, for example, the number of bedrooms in the third house

# What is the form of our hypothesis?

- Previously our hypothesis took the following form.
  - $h_{\theta}(x) = \theta_0 + \theta_1 x$ 
    - Here we have two parameters ( $\theta_1$  and  $\theta_2$ ) determined by our cost function
    - One variable  $x$
- Now we have multiple features
  - $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$
- Let us take  $x_0 = 1$  for convenience of notation
  - $h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$
  - $h_{\theta}(x) = \theta^T X$

# Gradient descent for multiple variables

- Our cost function is

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Gradient descent

Repeat {  
→  $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$   
} (simultaneously update for every  $j = 0, \dots, n$ )

# Gradient decent for multiple variables

- When  $n = 1$

Repeat {

$$\theta_0 := \theta_0 - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})}_{\frac{\partial}{\partial \theta_0} J(\theta)}$$
$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

(simultaneously update  $\theta_0, \theta_1$ ) }

- When  $n \geq 1$

New algorithm ( $n \geq 1$ ):

Repeat {

$\downarrow \frac{\partial}{\partial \theta_j} J(\theta)$

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$


(simultaneously update  $\theta_j$  for  
 $j = 0, \dots, n$ ) }



# Gradient Decent in practice: 1

## Feature Scaling

- Range difference of features
  - $x_1$  = size (0 - 2000 feet)
  - $x_2$  = number of bedrooms (1-5)
  - Means the contours generated if we plot  $\theta_1$  vs.  $\theta_2$  give a very tall and thin shape due to the huge range difference
- Running gradient descent on this kind of cost function can take a long time to find the global minimum
- Feature scaling
  - If you have a problem with multiple features
  - You should make sure those features have a similar scale
    - Means gradient descent will converge more quickly




# Some feature scaling methods

- Max/Min Scaling
- Mean Normalization
- Z-score Scaling



# Learning Rate $\alpha$

- Focus on the learning rate ( $\alpha$ )
  - How to choose  $\alpha$ ?
- Make sure gradient descent is working
- Plot  $\min J(\theta)$  vs. no of iterations
  - (i.e. plotting  $J(\theta)$  over the course of gradient descent)
- If gradient descent is working then  $J(\theta)$  should decrease after every iteration



# Learning Rate $\alpha$

- Checking its working If you plot  $J(\theta)$  vs. iterations and see the value is increasing - means you probably need a smaller  $\alpha$ 
  - Cause is because your minimizing a function which looks like this



- But you overshoot, so reduce learning rate so you actually reach the minimum (green line)

•

# Vector Notation

Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
$x_1$	$x_2$	$x_3$	$x_4$	$y$
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178

$$\bullet \quad X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}_{m \times (n+1)} \quad Y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}_{m \times 1}$$

$$\bullet \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}_{(n+1) \times 1} \quad h_{\theta}(x) = x_0\theta_0 + \cdots + x_n\theta_n = X\theta$$

$$\bullet \quad J = \frac{1}{2m} (X\theta - Y)^T_{1 \times m} (X\theta - Y)_{m \times 1}$$

# Vectorization

New algorithm ( $n \geq 1$ ):

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update  $\theta_j$  for  $j = 0, \dots, n$ ) }

$$\bullet \quad \frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) \cdot x_j^i$$

$$\bullet \quad = \frac{1}{m} \sum_{i=1}^m ((x^i \theta) - y^i) \cdot x_j^i$$

$$\bullet \quad \Theta_{(n+1) \times 1} = \Theta_{(n+1) \times 1} - \frac{\alpha}{m} (X^T_{(n+1) \times m} (X\theta - Y)_{m \times 1})$$

# Closed form solution

- Given  $X$  and  $Y$ , our aim is to find  $\theta$  so that  $Y = X\theta$ .
- $Y = X\theta$ .
- $X^{-1}Y = \theta$
- $X^{-1}(X^{T^{-1}}X^T)Y = \theta$
- $(X^TX)^{-1}X^TY = \theta$

# Why not closed form?

- The problem with this operation is the time complexity of calculating the inverse of a  $n \times n$  matrix which is  $O(n^3)$  and as  $n$  increases it can take a very long time to finish.
- When  $n$  is low ( $n < 1000$  or  $n < 10000$ ) you can think of normal equations as the better option for calculation theta, however for greater values **Gradient Descent** is much more faster, so the only reason is the time.
- Closed form works better when the input size is smaller. No need to choose  $\alpha$  and no need to iterate.



# Batch, Stochastic and Mini-Batch Gradient Descent



- In Gradient Descent or Batch Gradient Descent, we use the whole training data for updating theta (or,  $w$ ).

$$w = w - \alpha \nabla_w J(w) \quad (6)$$

- In Stochastic Gradient Descent, we use only single training example for updating theta.

$$w = w - \alpha \nabla_w J(x^i, y^i; w) \quad (7)$$

- Mini-batch Gradient Descent lies in between of these two extremes, in which we can use a mini-batch (small portion) of training data for updating theta.

$$w = w - \alpha \nabla_w J(x^{\{i:i+b\}}, y^{\{i:i+b\}}; w) \quad (7)$$

# References

- 1) Bishop, Christopher M. "Pattern recognition and machine learning, 2006." Springer 60.1 (2012): 78-78.
- 2) <http://www.holehouse.org/mlclass/>
- 3) [https://www.mathsisfun.com/equation of line.htm](https://www.mathsisfun.com/equation_of_line.htm)  
|
- 4) <https://towardsdatascience.com/gradient-descent-algorithm-and-its-variants-10f652806a3>